

# Data Analytics Project: Phishing Website Classification using Machine Learning.

**Course:** Data Anallystics and (Data Driven Decision)

**Objective:** Apply descriptive and predictive analytics to classify websites as legitimate, suspicious, or phishy using the UCI Website Phishing dataset.

## 1. Dataset Description

This dataset is designed for the classification task of identifying whether a given website is legitimate potentially phishing or Suspicious. It contains a total of 1,353 website instances, each characterized by 9 numerical features and a binary target variable indicating the classification `Result`.

- Format: ARFF
- 1353 instances (websites)
- 9 features + 1 target ( `Result` )
- Classes:
  - `-1` : Phishing
  - `0` : Suspicious
  - `1` : Legitimate

```
In [1]: import arff
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
```

### 1.1 Loading and Previewing the DataSet

```
In [3]: # Loading the arff datafile
with open('PhishingData.arff', 'r') as f:
    dataset = arff.load(f)

# converting the dataset to dataframe
df = pd.DataFrame(dataset['data'], columns=[attr[0] for attr in dataset['attributes']])

# preview the first 5 rows
display(df.head())
```

	SFH	popUpWidnow	SSLfinal_State	Request_URL	URL_of_Anchor	web_traffic	URL_Length	age_of_domain	having_IP_Address	Result
0	1	-1	1	-1	-1	1	1	1	0	0
1	-1	-1	-1	-1	-1	0	1	1	1	1
2	1	-1	0	0	-1	0	-1	1	0	1
3	1	0	1	-1	-1	0	1	1	0	0
4	-1	-1	1	-1	0	0	-1	1	0	1

### 1.2 checking the shape of the dataset

```
In [5]: # checking the shape of the dataset
display("Dataset shape", df.shape)

'Dataset shape'
(1353, 10)

In [6]: #column names and datatypes
print("\nColumn data type:")
print(df.dtypes)
```

Column data type:  
SFH object  
popUpWidnow object  
SSLfinal\_State object  
Request\_URL object  
URL\_of\_Anchor object  
web\_traffic object  
URL\_Length object  
age\_of\_domain object  
having\_IP\_Address object  
Result object  
dtype: object

```
In [7]: #summary statistics
print("\nSummary Statistics:")
display(df.describe())
```

Summary Statistics:

	SFH	popUpWidnow	SSLfinal_State	Request_URL	URL_of_Anchor	web_traffic	URL_Length	age_of_domain	having_IP_Address	Result
count	1353	1353	1353	1353	1353	1353	1353	1353	1353	1353
unique	3	3	3	3	3	3	3	2	2	2
top	1	0	1	-1	-1	0	0	1	0	0
freq	767	639	751	617	610	473	563	825	1198	1198

```
In [8]: #checking class distribution in the target variable
print("\nTarget class distribution:")
print(df['Result'].value_counts())
```

Target class distribution:  
Result  
-1 702  
1 548  
0 103  
Name: count, dtype: int64

1.3 Dataset Summary

- Shape: 1353 rows x 10 coulmns
- Target columns: ( Result )
- Unique target values: -1 , 0 , 1 .

This means that the dataset is likely Multi-class classification, where:

- -1 = Phishy (or malicious)
- 0 = Suspicious or neutral
- 1 = Legitimate.

2 Data cleaing

Checking for missing values, Duplicates or any data type issues

```
In [11]: #checking for missing values
print("Missing values per column")
print(df.isnull().sum())
```

Missing values per column  
SFH 0  
popUpWidnow 0  
SSLfinal\_State 0  
Request\_URL 0  
URL\_of\_Anchor 0  
web\_traffic 0  
URL\_Length 0  
age\_of\_domain 0  
having\_IP\_Address 0  
Result 0  
dtype: int64

```
In [12]: #checking for duplicates
print("\nNumber of duplicate rows:", df.duplicated().sum())
```

Number of duplicate rows: 629

2.1 Data cleaning summary

- No missing values : There is no need to fill or drop rows due to missing data.
- 629 duplicate rows : That is approximately 46% of the dataset!

We decided to keep duplicates becasue:

- Each row is a valid, independent observation, even if values match.
- The dataset behavior-based phishing indicators, it’s possible that,Some websites exhibit identical behavior and hence should not be dropped.

```
In [14]: #Keep the Original Dataset shape
print ("We keep the Original shape of the Dataset:", df.shape)
```

We keep the Original shape of the Dataset: (1353, 10)

### 3 Exploratory Data Analysis (EDA)

- Understanding the data visually and statistically, Identify trends. patterns, and feature relationships. Get insights into the Result class distribution.

#### 3.1 Basic analysis

```
In [17]: # Ensuring that all the values are numeric
df = df.apply(pd.to_numeric)

#getting baic statistic for all features
basic_stats =df.describe()
display(basic_stats)
```

	SFH	popUpWidnow	SSLfinal_State	Request_URL	URL_of_Anchor	web_traffic	URL_Length	age_of_domain	having_IP_Addre
count	1353.000000	1353.000000	1353.000000	1353.000000	1353.000000	1353.000000	1353.000000	1353.000000	1353.0000
mean	0.237990	-0.258684	0.327421	-0.223208	-0.025129	0.000000	-0.053215	0.219512	0.1145
std	0.916389	0.679072	0.822193	0.799682	0.936262	0.806776	0.762552	0.975970	0.3186
min	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	0.0000
25%	-1.000000	-1.000000	0.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	0.0000
50%	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.0000
75%	1.000000	0.000000	1.000000	0.000000	1.000000	1.000000	1.000000	1.000000	0.0000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.0000

#### 3.2 standard deviation analysis

```
In [19]: print("Standard deviation per feature:\n")
print(df.std())
```

Standard deviation per feature:

```
SFH          0.916389
popUpWidnow  0.679072
SSLfinal_State  0.822193
Request_URL   0.799682
URL_of_Anchor 0.936262
web_traffic   0.806776
URL_Length    0.762552
age_of_domain 0.975970
having_IP_Address 0.318608
Result        0.954773
dtype: float64
```

#### Basic Statistical Analysis

We performed basic statistical analysis using `df.describe()` . This provided insights into the distribution of each feature including mean, min, max, and standard devatio##


#### Standard Deviation Analysis

The standard deviation tells us how much the values for each feature vary from the mean. A higher standard deviation indicates greater variability across websites, while a lower standard deviation suggests uniform behavior for that feature.

```
In [21]: #set seaborn style
# Set Seaborn style
sns.set(style='whitegrid')

# 1. Target class distribution
plt.figure(figsize=(6,4))
sns.countplot(x='Result', data=df, palette='Set2')
```

```
plt.title('Target Class Distribution')
plt.xlabel('Result')
plt.ylabel('Count')
plt.show()
```



### 3.1 class distribution Analysis


- -1 = Phishing with 700 count is the Majority class .
- 1 = Legitimate with 550 count is the Close Second .
- 0 = Suspicious with 100 count is the Minor Class(rare cases) .

This is slightly imbalanced dataset , especially for class 0 .

### 3.4 Correlation Matrix

Most features are weakly correlated, meaning they provide independent information, which is useful for machine learning models

```
In [33]: # 2. Heatmap to visualize correlation between features
plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Feature Correlation Matrix')
plt.show()
```



### 3.5 PCA

In the case of our Dataset , PCA is best use for visualization only, not as a requirement before modelling, because the features are already informative and interpretable. The features like( SFH , SSLfinal\_State , etc.) are:

- Discrete, categorical-like values (-1, 0, 1)
- Already engineered by experts to reflect phishing behaviors


These features have clear meanings and are already effective as inputs.

```
In [38]: # Prepare data
X = df.drop('Result', axis=1)
y = df['Result']

# Apply PCA to reduce to 2D
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Create a DataFrame with PCA results
pca_df = pd.DataFrame(data=X_pca, columns=['PC1', 'PC2'])
pca_df['Result'] = y.values

# Plot
plt.figure(figsize=(8,6))
sns.scatterplot(data=pca_df, x='PC1', y='PC2', hue='Result', palette='Set2')
plt.title("PCA - 2D Projection of Feature Space")
plt.grid(True)
plt.show()
```

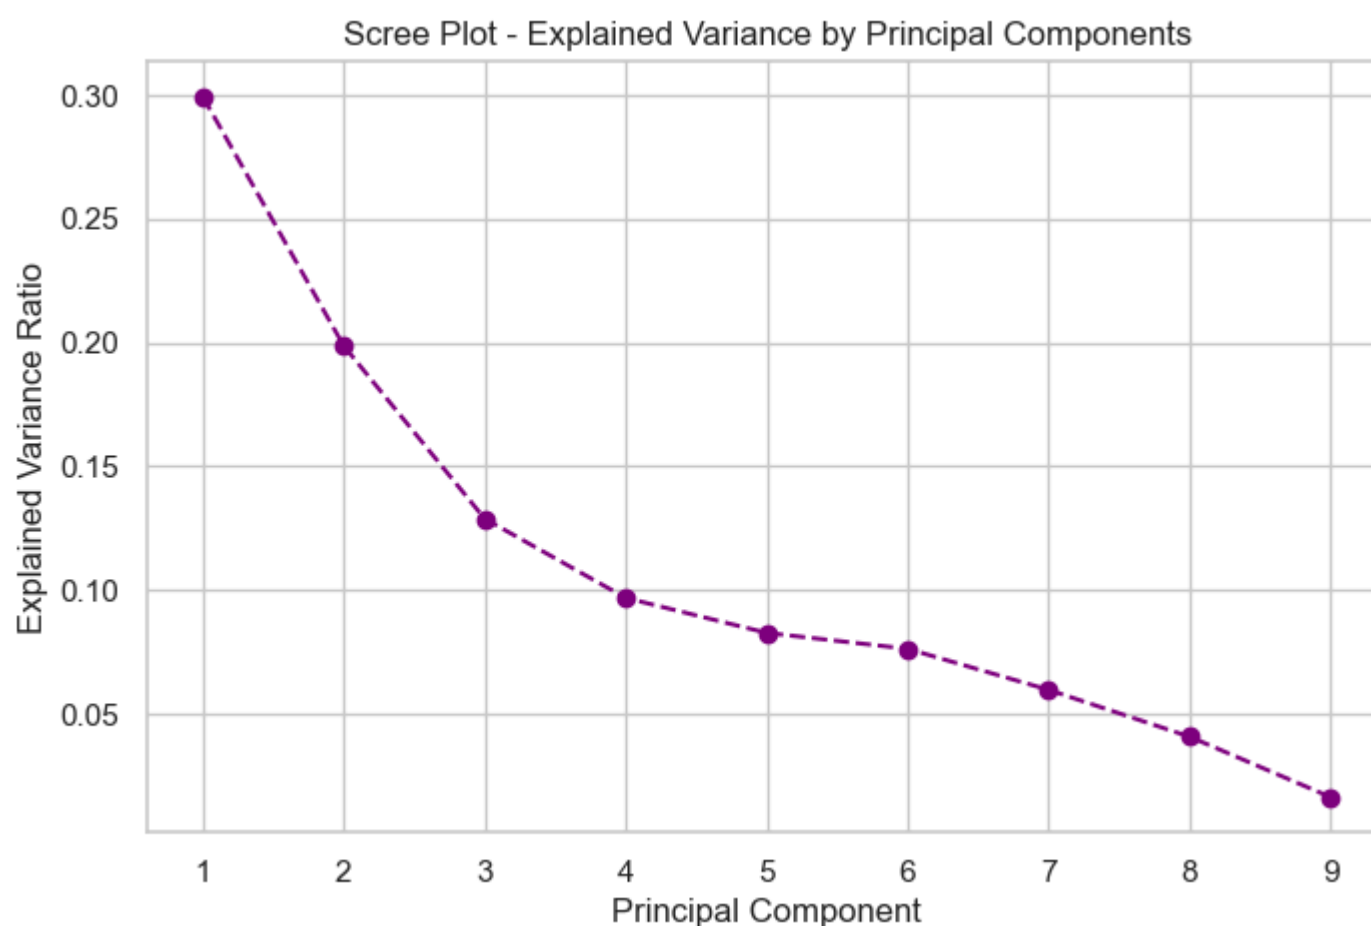


### 3.6 scree plot

```
In [41]: # Prepare feature matrix (excluding the target)
X = df.drop('Result', axis=1)

# Apply PCA without limiting the number of components
pca = PCA()
pca.fit(X)

# Plot the explained variance for each principal component
plt.figure(figsize=(8, 5))
plt.plot(
    range(1, len(pca.explained_variance_ratio_) + 1),
    pca.explained_variance_ratio_,
    marker='o',
    linestyle='--',
    color='purple'
)
plt.title('Scree Plot - Explained Variance by Principal Components')
plt.xlabel('Principal Component')
plt.ylabel('Explained Variance Ratio')
plt.xticks(range(1, X.shape[1] + 1))
plt.grid(True)
plt.show()
```



## 4 Supervised ML model

Base on our dataset: We are working on a **Supervised Learning Task**

- Data are lebeled: the **Result** column (values **-1 , 0 , 1** )
- Labels are already known in advance.

The **goal** is to use the labeled data to train a model to predict **Results** based on the feafures.

### Main Analysis

- split the data into training and testing sets
- train a model
- evalute how well it performs.

#### 4.1 Data splitting

- x = Features(everything except Results)
- y = Target (Result column)

```
In [47]: # Features and Target
X = df.drop('Result', axis=1)
y = df['Result']

# Splitting the dataset (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42) #We use stratify=y to ensure all 3 classes are proportionally represented

# Check the shape of the splits
print("Training set:", X_train.shape)
print("Testing set:", X_test.shape)
```

Training set: (1082, 9)

Testing set: (271, 9)

#### 4.2 Training the Decision Tree

```
In [50]: # Initialize and train the classifier
clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)

# Predict on test data
y_pred = clf.predict(X_test)
```

#### 4.3 Evaluateing the model

```
In [53]: # Confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[128  3  9]
 [ 4 13  4]
 [ 10  2 98]]
```

Classification Report:

	precision	recall	f1-score	support
-1	0.90	0.91	0.91	140
0	0.72	0.62	0.67	21
1	0.88	0.89	0.89	110
accuracy			0.88	271
macro avg	0.84	0.81	0.82	271
weighted avg	0.88	0.88	0.88	271

### Try with random forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

# Initialize and train the Random Forest classifier
rf_clf = RandomForestClassifier(random_state=42)
```

```
rf_clf.fit(X_train, y_train)

# Predict on test data
rf_pred = rf_clf.predict(X_test)

# Evaluation
print("Confusion Matrix:")
print(confusion_matrix(y_test, rf_pred))

print("\nClassification Report:")
print(classification_report(y_test, rf_pred))
```

Confusion Matrix:

[[126	2	12]
[ 5	12	4]
[ 8	2	100]]

Classification Report:

	precision	recall	f1-score	support
-1	0.91	0.90	0.90	140
0	0.75	0.57	0.65	21
1	0.86	0.91	0.88	110
accuracy			0.88	271
macro avg	0.84	0.79	0.81	271
weighted avg	0.88	0.88	0.88	271

## Model Comparison – Decision Tree vs Random Forest

Both the Decision Tree and Random Forest classifiers achieved an overall accuracy of 88%, showing that each model is effective at detecting phishing websites. However, a deeper look at the precision, recall, and F1-scores reveals some key differences in performance.

Decision Tree:

- Slightly better recall and F1-score for the suspicious class (0).
- Strong performance across phishing (-1) and legitimate (1) classes.

Random Forest:

- Slightly higher precision on the suspicious class (0), but lower recall compared to Decision Tree.
- More stable across classes due to being an ensemble of multiple trees.
- Less prone to overfitting and generally more robust, especially on unseen data.

## SUMMARY FOR THE LEARNING MODEL

Why We Used Supervised Learning.

We chose supervised learning because our problem involved predicting known classes, i.e. phishing (-1), suspicious (0), and legitimate (1) based on labeled historical data.

Reasons.

- The dataset includes a target variable (Result), which clearly labels each website.
- Our goal was to learn patterns from these labeled examples to predict future website behavior.
- Models like Decision Trees and Random Forests are well-suited for this kind of task.

Why we did not used Unsupervised Learning.

- Unsupervised learning is used when there's no labeled output for example, clustering websites based on similar behaviors without knowing if they're phishing or not.
- It's useful for exploration, but it won't give exact classifications like "this is a phishing site."

Final Thought: Because we had clear labels and a well-defined classification task, supervised learning was the most appropriate and effective choice for detecting phishing websites.

## Conclusions

**Descriptive analytics** We began by exploring the dataset using descriptive analytics. This helped us to:

- Understand the distribution of the classes (phishing -1 , suspicious 0 , legitimate 1 )
- Analyze the most common behaviors of websites, such as the use of SSL certificates, IP addresses, form handler etc.
- Revealed class imbalanced, with more legitimate and phishy sites.
- Generate summary statistics and visualizations (e.g., bar charts, correlation heatmaps).



Therefore; Descriptive analytics gave us the insight we needed to understand the dataset, detect patterns, and make informed decisions about how to clean the data and what models to use.

### **Predictive analytics**

Predictive analytics enabled us to build models that could automatically classify websites. We, therefore; applied:

- Decision Tree and Random Forest models, both achieving an accuracy of 88%. Using both the models achieved reasonable classification performance.
- Feature importance analysis, which revealed that certain website characteristics (e.g., SSL use, URL structure) are strong predictors of phishing behavior.
- Feature importance analysis highlights which website characteristics are most predictive.
- Improvements: Try more models (Logistics regression,hyperparameter tuning, SVM), hyperparameter tuning, and cross-validation.